



Bookmarks

- ▶ Week 1
- ▶ Week 2
- ▶ Week 3
- ▶ Week 4
- ▶ Week 5

▼ Week 6

Lecture 10:
Reinforcement
Learning (edited)

Lecture 10:
Reinforcement
Learning (live)

Lecture 11:
Reinforcement
Learning II (edited)

Lecture 11:
Reinforcement
Learning II (live)

Homework 5:
Reinforcement
Learning
Homework

Project 3:
Reinforcement
Learning
Project 3

Midterm 1
Preparation

Processing math: 40%

Week 6 > Project 3: Reinforcement Learning >
p3_rl_q8_approximate_qlearning_and_features

Bookmark

Question 8 (3 points): Approximate Q-Learning

Implement an approximate Q-learning agent that learns weights for features of states, where many states might share the same features. Write your implementation in `ApproximateQAgent` class in `qlearningAgents.py`, which is a subclass of `PacmanQAgent`.

Note: Approximate Q-learning assumes the existence of a feature function $f(s,a)$ over state and action pairs, which yields a vector $f_1(s,a) \dots f_i(s,a) \dots f_n(s,a)$ of feature values. We provide feature functions for you in `featureExtractors.py`. Feature vectors are `util.Counter` (like a dictionary) objects containing the non-zero pairs of features and values; all omitted features have value zero.

The approximate Q-function takes the following form

$$Q(s, a) = \sum_{i=1}^n f_i(s, a) w_i$$

where each weight w_i is associated with a particular feature $f_i(s,a)$. In your code, you should implement the weight vector as a dictionary mapping features (which the feature extractors will return) to weight values. You will update your weight vectors similarly to how you updated Q-values:

$$w_i \leftarrow w_i + \alpha \cdot \text{difference} \cdot f_i(s, a)$$

$$\text{difference} = (r + \gamma \max_{a'} Q(s, a') - Q(s, a))$$

Note that the difference term is the same as in normal Q-learning, and r is the experienced reward.

By default, `ApproximateQAgent` uses the `IdentityExtractor`, which assigns a single feature to every $(\text{state}, \text{action})$ pair. With this feature extractor, your approximate Q-learning agent should work identically to `PacmanQAgent`. You can test this with the following command:

▶ Week 7

```
python pacman.py -p ApproximateQAgent -x 2000 -n 2010 -l
smallGrid
```

▶ Week 8

Important: ApproximateQAgent is a subclass of QLearningAgent, and it therefore shares several methods like `getAction`. Make sure that your methods in QLearningAgent call `getQValue` instead of accessing Q-values directly, so that when you override `getQValue` in your approximate agent, the new approximate q-values are used to compute actions.

▶ Week 9

▶ Week 10

▶ Week 11

Once you're confident that your approximate learner works correctly with the identity features, run your approximate Q-learning agent with our custom feature extractor, which can learn to win with ease:

▶ Week 12

▶ Week 13

```
python pacman.py -p ApproximateQAgent -a
extractor=SimpleExtractor -x 50 -n 60 -l mediumGrid
```

▶ Week 14

Even much larger layouts should be no problem for your ApproximateQAgent. (*warning:* this may take a few minutes to train)

```
python pacman.py -p ApproximateQAgent -a
extractor=SimpleExtractor -x 50 -n 60 -l mediumClassic
```

If you have no errors, your approximate Q-learning agent should win almost every time with these simple features, even with only 50 training games.

Grading: We will run your approximate Q-learning agent and check that it learns the same Q-values and feature weights as our reference implementation when each is presented with the same set of examples. To grade your implementation, run the autograder:

```
python autograder.py -q q8
```

Congratulations! You have a learning Pacman agent!

© All Rights Reserved



Processing math: 40%

© edX Inc. All rights reserved except where noted. EdX, Open edX and the edX and Open EdX logos are registered trademarks or trademarks of edX Inc.

POWERED BY
OPENedX

