



Bookmarks



Bookmark

▶ Week 1

▶ Week 2

▶ Week 3

▶ Week 4

▶ Week 5


▼ Week 6


Lecture 10:
Reinforcement
Learning (edited)

Lecture 10:
Reinforcement
Learning (live)

Lecture 11:
Reinforcement
Learning II (edited)

Lecture 11:
Reinforcement
Learning II (live)

Homework 5:
Reinforcement
Learning
Homework 

**Project 3:
Reinforcement
Learning**
Project 3 

Midterm 1
Preparation

Week 6 > Project 3: Reinforcement Learning > p3_rl_welcome

MDPs

To get started, run Gridworld in manual control mode, which uses the arrow keys:

```
python gridworld.py -m
```

You will see the two-exit layout from class. The blue dot is the agent. Note that when you press *up*, the agent only actually moves north 80% of the time. Such is the life of a Gridworld agent!

You can control many aspects of the simulation. A full list of options is available by running:

```
python gridworld.py -h
```

The default agent moves randomly

```
python gridworld.py -g MazeGrid
```

You should see the random agent bounce around the grid until it happens upon an exit. Not the finest hour for an AI agent.

Note: The Gridworld MDP is such that you first must enter a pre-terminal state (the double boxes shown in the GUI) and then take the special 'exit' action before the episode actually ends (in the true terminal state called `TERMINAL_STATE`, which is not shown in the GUI). If you run an episode manually, your total return may be less than you expected, due to the discount rate (`-d` to change; 0.9 by default).

Look at the console output that accompanies the graphical output (or use `-t` for all text). You will be told about each transition the agent experiences (to turn this off, use `-q`).

As in Pacman, positions are represented by (x, y) Cartesian coordinates and any arrays are indexed by $[x][y]$, with 'north' being the direction of increasing y , etc. By default, most transitions will receive a reward of zero,

- ▶ Week 7
- ▶ Week 8
- ▶ Week 9
- ▶ Week 10
- ▶ Week 11
- ▶ Week 12
- ▶ Week 13
- ▶ Week 14

though you can change this with the living reward option (-r).

© All Rights Reserved



© edX Inc. All rights reserved except where noted. EdX, Open edX and the edX and Open EdX logos are registered trademarks or trademarks of edX Inc.

POWERED BY
OPENedX

